

Final Team Project - Airbnb Boston Data

Step I: Data Preparation & Exploration

Read data into your local environment

Hide

```
df <- read.csv("metad699_train.csv")  
View(df)  
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Hide

```
library(tidyr)  
library(tidyverse)
```

```
[37m— [1mAttaching packages [22m ————— tidyvers
e 1.3.0 — [39m
[37m [32m✓ [37m [34mggplot2 [37m 3.2.1      [32m✓ [37m [34mpurrr [37m 0.3.3
[32m✓ [37m [34mtibble [37m 2.1.3      [32m✓ [37m [34mstringr [37m 1.4.0
[32m✓ [37m [34mreadr [37m 1.3.1      [32m✓ [37m [34mforcats [37m 0.4.0 [39m
[37m— [1mConflicts [22m ————— tidyverse_conf
licts() —
[31mx [37m [34mdplyr [37m:: [32mfilter() [37m masks [34mstats [37m::filter()
[31mx [37m [34mdplyr [37m:: [32mlag() [37m masks [34mstats [37m::lag() [39m
```

Hide

```
library(ggplot2)
library(caret)
```

Loading required package: lattice

Attaching package: 'caret'

The following object is masked from 'package:purrr':

```
lift
```

Hide

```
library(plyr)
```

```
-----  
You have loaded plyr after dplyr - this is likely to cause problems.  
If you need functions from both plyr and dplyr, please load plyr first, then dplyr:  
library(plyr); library(dplyr)  
-----
```

Attaching package: 'plyr'

The following object is masked from 'package:purrr':

compact

The following objects are masked from 'package:dplyr':

arrange, count, desc, failwith, id, mutate, rename, summarise, summarize

[Hide](#)

```
library(forecast)
```

Registered S3 method overwritten by 'quantmod':

```
method      from  
as.zoo.data.frame zoo
```

This is forecast 8.11

Want to meet other forecasters? Join the International Institute of Forecasters:
<http://forecasters.org/>

[Hide](#)

```
boston <- filter(df, city=="Boston")
```

I. Missing Values

[Code](#)

```
[1] "We first used the anyNA function to determine if we had any missing data. Upon filtering the data with our selected city of Boston, we were able to observe a significant number of blank cells and missing values. We decided to find and use the median values as replacements and replaced blank spaces with NA's. Overall, we ended up preserving our full filtered data with 3468 observations and 29 variables. Upon preserving all our data for numerical categories, we deleted unnecessary columns and further cleaned our data frame where no missing NA's were present in colSums. We re-named our finalized copy as 'boston1'."
```

Hide

```
anyNA(boston)
```

```
[1] TRUE
```

Hide

```
# Explore missing values
View(boston)
colSums(is.na(boston))
```

id	log_price	property_type	room_type
0	0	0	0
amenities	accommodates	bathrooms	bed_type
0	0	6	0
cancellation_policy	cleaning_fee	city	description
0	0	0	0
first_review	host_has_profile_pic	host_identity_verified	host_response_rate
0	0	0	0
host_since	instant_bookable	last_review	latitude
0	0	0	0
longitude	name	neighbourhood	number_of_reviews
0	0	0	0
review_scores_rating	thumbnail_url	zipcode	bedrooms
648	0	0	3
beds			
2			

Hide

```
# Explore median values for missing column without factoring NA's  
median(boston$review_scores_rating, na.rm = TRUE)
```

```
[1] 96
```

[Hide](#)

```
median(boston$bathrooms, na.rm = TRUE)
```

```
[1] 1
```

[Hide](#)

```
median(boston$bedrooms, na.rm = TRUE)
```

```
[1] 1
```

[Hide](#)

```
median(boston$bathrooms, na.rm = TRUE)
```

```
[1] 1
```

[Hide](#)

```
median(boston$bedrooms, na.rm = TRUE)
```

```
[1] 1
```

[Hide](#)

```
median(boston$beds, na.rm=TRUE)
```

```
[1] 1
```

Hide

```
# replace all NA's
boston[boston==" "] <-NA
# replace all NA's with median value
boston$review_scores_rating[is.na(boston$review_scores_rating)] <- median(boston$review_scores_rating, na.rm = TRUE)
boston$host_response_rate <- as.numeric(sub("%","",boston$host_response_rate))/100
boston$host_response_rate[is.na(boston$host_response_rate)] <- median(boston$host_response_rate, na.rm=TRUE)
boston$beds[is.na(boston$beds)] <- median(boston$beds, na.rm = T)
boston$bathrooms[is.na(boston$bathrooms)] <- median(boston$bathrooms, na.rm=T)
boston$bedrooms[is.na(boston$bedrooms)] <- median(boston$bedrooms, na.rm=T)
colSums(is.na(boston))
```

id	log_price	property_type	room_type
0	0	0	0
amenities	accommodates	bathrooms	bed_type
0	0	0	0
cancellation_policy	cleaning_fee	city	description
0	0	0	0
first_review	host_has_profile_pic	host_identity_verified	host_response_rate
621	0	0	0
host_since	instant_bookable	last_review	latitude
0	0	621	0
longitude	name	neighbourhood	number_of_reviews
0	0	0	0
review_scores_rating	thumbnail_url	zipcode	bedrooms
0	134	26	0
beds			
0			

Hide

```
# delete unnecessary information
boston1 <- boston[-c(13, 11, 19, 26, 27)]
View(boston1)
colSums(is.na(boston1))
```

id	log_price	property_type	room_type
0	0	0	0
amenities	accommodates	bathrooms	bed_type
0	0	0	0
cancellation_policy	cleaning_fee	description	host_has_profile_pic
0	0	0	0
host_identity_verified	host_response_rate	host_since	instant_bookable
0	0	0	0
latitude	longitude	name	neighbourhood
0	0	0	0
number_of_reviews	review_scores_rating	bedrooms	beds
0	0	0	0

II. Summary Statistics

Code

```
[1] "In our selected data frame, we wanted to observe the true nightly price from log_price to understand the true dollar format. Upon running the summary of boston1 we noticed the following observations for the selected variables:\n\n\t\tReview score ratings: Out of 100 being the highest score for review ratings from \t\t\tcustomers, 20 was the lowest. The median was 96 and the mean was 94.05 which makes the distribution of review ratings negatively skewed. When we further observe the standard deviation of 7.327312, we can conclude that our review ratings are very close to the mean.\n\n\t\tLog price/Nightly price: We know that log price and nightly price are practically the same. The max rental price for an Airbnb in Boston is $1,400 a night while the lowest is $17. The median price is $136 while the mean is $165.50 which means that \tmedian is low and mean is high. In contrast to log price form, the median is higher while the mean is lower. This means that the dollar form of our nightly price data is \tpositively skewed while in log form, the price is negatively skewed. A possibility \tfor why this happens is the normalization of data where the price in normal format can be far spread out. This makes sense if we were to observe the standard deviation for both prices. In log format, the standard deviation is close to 0 which \t\tmeans that all the data point are close to its mean while the nightly price has a standard deviation farther apart from mean. \n"
```

Hide

```
boston1$nightly_price <- exp(boston1$log_price) # nightly price conversion from log
# Selected summary of statistics
summary(boston1$review_scores_rating)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
20.00	92.00	96.00	94.05	98.00	100.00

Hide

```
summary(boston1$log_price)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2.833	4.382	4.913	4.884	5.298	7.244

Hide

```
summary(boston1$nightly_price)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
17.0	80.0	136.0	165.6	200.0	1400.0

Hide

```
sd(boston1$review_scores_rating)
```

```
[1] 7.327312
```

Hide

```
sd(boston1$log_price)
```

```
[1] 0.6646924
```

Hide

```
sd(boston1$nightly_price)
```

```
[1] 128.8892
```



III. Visualization

[Code](#)


```
[1] "For histogram and scatter plot we saw that our scores rating remained the same. Using the Histogram, we can see that the majority of our data is skewed to the right where the scores tend to be above 75%. Only few outliers exist where the review scores are less than 50%. The same can be confirmed with the Scatter Plot where the bulk of the reviews are rated at above 80%. This indicates that the reviews for Airbnb rentals in Boston are positive\n\nWe further expanded the exploration of our data by choosing neighborhood and log price as variables. If there is one thing that influences someone to book a room, house, or apartment, it is the price you pay to book your rental. In our bar plot and violin plot, we can observe the log price per neighborhood. In the given bar plot, we see that the Allston-Brighton carries the bulk of the rentals based on prices. This means that those who rent through Airbnb, would rent the most in that area with the given price range. The Violin Plot gives a better indication of where that range lies. Based on the distribution of prices for Allston-Brighton, the bulk for log price lies slightly above 3 and up to about 5. We can conclude that the area is a frequent Airbnb hotspot due to its lower prices.\n\nWhen we examine our boxplot, in reference to review scores ratings and cancellation policies, the flexible a rental is, the higher the review will be. We can see that as the box plot becomes larger as the cancellation policy becomes more strict.\n"
```

[Hide](#)

```
ggplot(boston1, aes(x=review_scores_rating)) +  
  geom_histogram(color="darkblue", fill="lightblue") +  
  labs(title="Histogram of Review Scores Rating") +  
  theme_classic()
```

[Hide](#)

```
ggplot(boston1, aes(x=review_scores_rating, y=number_of_reviews, color=review_scores_rating)) +  
  geom_point() +  
  labs(title="Scatter Plot of Review Scores Rating") +  
  theme_classic()
```

[Hide](#)

```
ggplot(boston1, aes(x=neighbourhood, y=log_price, fill=neighbourhood)) +  
  geom_bar(stat = "identity") +  
  labs(title="Bar Plot of Price per neighborhood") +  
  theme(axis.text.x = element_text(angle=45, hjust=1))
```



Hide

```
install.packages("Hmisc")
```

```
Installing package into '/Users/josemartinez/Library/R/3.6/library'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/macosx/el-capitan/contrib/3.6/Hmisc_4.4-0.tgz'
Content type 'application/x-gzip' length 3146788 bytes (3.0 MB)
=====
downloaded 3.0 MB
```

The downloaded binary packages are in
/var/folders/6v/wsr694r57n9dfsdxfthdysbh0000gn/T//Rtmp8P9CxE/downloaded_packages

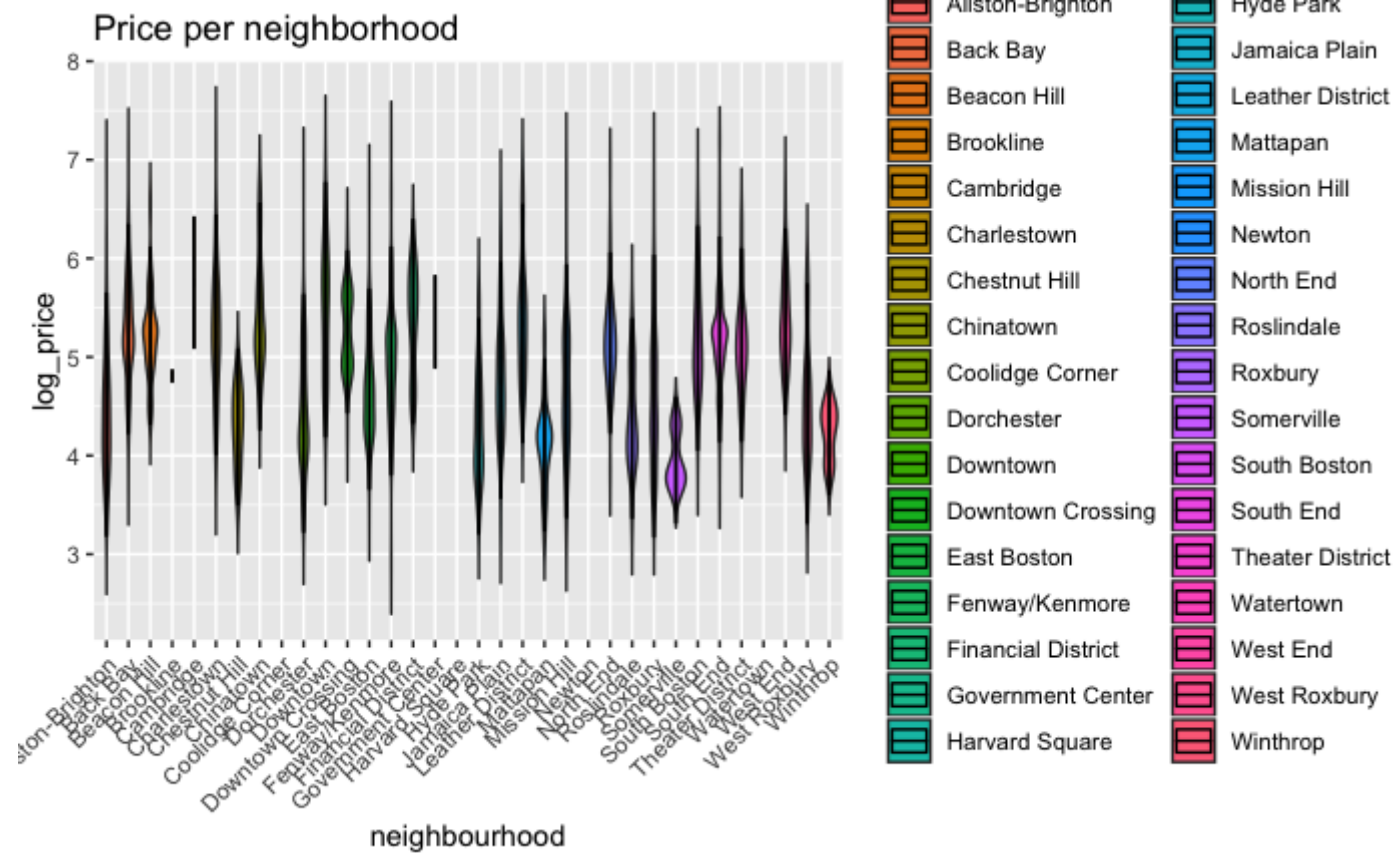
[Hide](#)

```
Violin <- ggplot(boston1, aes(neighbourhood, y=log_price, fill=neighbourhood)) +  
  geom_violin(trim=FALSE) +  
  stat_summary(fun.data="mean_sdl", mult=1, geom="crossbar", width=0.04 ) +  
  labs(title = "Price per neighborhood") + theme(axis.text.x = element_text(angle=45, hjust=1))
```


Ignoring unknown parameters: mult

[Hide](#)

Violin




```
ggplot(boston1, aes(cancellation_policy, review_scores_rating)) + geom_boxplot(fill="plum") +
  labs(title="Distribution of Review Scores Rating across various Cancellation ")
```



Step II: Prediction

[Code](#)

```
[1] "We began by setting up a correlation table and looking at the variables that were heavily correlated with each other. Due to the numerous amounts of variables that were provided in the boston1 dataframe, we knew that our best option was to select a multiple linear regression model to determine our prediction. We used the supply function to vector all columns and create a matrix. Based on the correlation table that was created without excluding any data, we saw that log_price and nightly_price were completely correlated with each other. We also noticed that beds, bedrooms, and accommodates were heavily correlated. Overall, we decided to remove id, nightly_price, beds, and bedrooms to prevent multicollinearity. The below heatmap, indicates that there is no issue of multicollinearity.\n\nWe further expanded the selection of our variables by using the backward elimination method in our multiple linear regression model. We used the 60/40 method to slice our data and train our model before validating it. Upon running the backward elimination, we saw that our recommended variables were narrowed down to 14 with an intercept of 1.20057241. If we were to determine our log_price for any given coefficient such accommodates as displayed in our regression summary, our regression formula would be as follows,\n\nlog_price = 1.200 + 0.0819 * accommodates\n\nAssume you want to accommodate for 3 people, the equation would be as follows,\n\nlog_price = 1.200 + 0.0819 * 3\nlog_price = 1.4457\n\nThe r-squared for our model is 0.5994. This means that close to 60% of our selected variables points would fit on the regression line. Our RMSE is 0.4444 which measures the difference between predicted values and actual values. The closer the number is to 0, the better. \n"
```

Hide

```
# normalize all data points
install.packages("reshape")
```

```
Installing package into '/Users/josemartinez/Library/R/3.6/library'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/macosx/el-capitan/contrib/3.6/reshape_0.8.8.tgz'
Content type 'application/x-gzip' length 172673 bytes (168 KB)
=====
downloaded 168 KB
```

```
The downloaded binary packages are in
/var/folders/6v/wsr694r57n9dfsdxfthdysbh0000gn/T//RtmplIEjsq/downloaded_packages
```

Hide

```
library(reshape)
```


Attaching package: 'reshape'

The following objects are masked from 'package:plyr':

rename, round_any

The following objects are masked from 'package:tidyr':


expand, smiths

The following object is masked from 'package:dplyr':

rename

[Hide](#)

```
bos <- boston1
bos <- bos[-c(1, 23, 24 ,25)]
must_convert <- sapply(bos, is.factor)
m2 <- sapply(bos[, must_convert], unclass)
bos <- cbind(bos[,!must_convert], m2)
table <- cor(bos)
melted_table <- melt(table)
# using a heatmap for slected variables
library(ggplot2)
ggplot(data = melted_table, aes(X1, X2, fill = value))+
  geom_tile(color = "white")+
  scale_fill_gradient2(low = "yellow", high = "red", mid = "white",
                      midpoint = 0, limit = c(-1,1), space = "Lab",
                      name="Correlation") +
  theme_minimal()+
  theme(axis.text.x = element_text(angle = 45, vjust = 1,
                                    size = 12, hjust = 1)) + coord_fixed()
```



Hide

468 * 0.60

[1] 280.8

Hide

3468 * 0.40

[1] 1387.2

Hide

```
Training <- slice(bos, 1:2080)
Validation <- slice(bos, 2081:3468)
# using the backward elimination method to further finalize our variables
boston_mlr <- lm(log_price~ ., data = Training)
step_boston_mlr <- step(boston_mlr, direction = "backward")
```

Start: AIC=-3562.42

```
log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + cancellation_policy +
  description + host_has_profile_pic + host_identity_verified +
  host_since + instant_bookable + name + neighbourhood
```

	Df	Sum of Sq	RSS	AIC
- description	1	0.000	367.69	-3564.4
- name	1	0.003	367.69	-3564.4
- cancellation_policy	1	0.062	367.75	-3564.1
- host_since	1	0.072	367.76	-3564.0
- host_has_profile_pic	1	0.155	367.84	-3563.5
- host_identity_verified	1	0.342	368.03	-3562.5
<none>			367.69	-3562.4
- property_type	1	0.577	368.26	-3561.2
- number_of_reviews	1	0.602	368.29	-3561.0
- host_response_rate	1	0.918	368.60	-3559.2
- bed_type	1	0.939	368.62	-3559.1
- neighbourhood	1	0.998	368.68	-3558.8
- review_scores_rating	1	1.676	369.36	-3555.0
- amenities	1	1.861	369.55	-3553.9
- instant_bookable	1	2.040	369.73	-3552.9
- cleaning_fee	1	2.123	369.81	-3552.4
- bathrooms	1	5.416	373.10	-3534.0
- longitude	1	8.399	376.08	-3517.4
- latitude	1	13.090	380.78	-3491.7
- accommodates	1	39.772	407.46	-3350.8
- room_type	1	119.723	487.41	-2978.1

Step: AIC=-3564.42

```
log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + cancellation_policy +
  host_has_profile_pic + host_identity_verified + host_since +
  instant_bookable + name + neighbourhood
```

	Df	Sum of Sq	RSS	AIC
- name	1	0.003	367.69	-3566.4
- cancellation_policy	1	0.062	367.75	-3566.1

```

- host_since      1      0.072 367.76 -3566.0
- host_has_profile_pic 1      0.155 367.84 -3565.5
- host_identity_verified 1      0.342 368.03 -3564.5
<none>                                367.69 -3564.4
- property_type   1      0.577 368.26 -3563.2
- number_of_reviews 1      0.602 368.29 -3563.0
- host_response_rate 1      0.919 368.60 -3561.2
- bed_type        1      0.939 368.62 -3561.1
- neighbourhood   1      1.001 368.69 -3560.8
- review_scores_rating 1      1.676 369.36 -3557.0
- amenities       1      1.864 369.55 -3555.9
- instant_bookable 1      2.045 369.73 -3554.9
- cleaning_fee    1      2.127 369.81 -3554.4
- bathrooms       1      5.417 373.10 -3536.0
- longitude       1      8.407 376.09 -3519.4
- latitude        1     13.097 380.78 -3493.6
- accommodates    1     39.810 407.50 -3352.6
- room_type       1    119.810 487.50 -2979.8

```

Step: AIC=-3566.4

```

log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + cancellation_policy +
  host_has_profile_pic + host_identity_verified + host_since +
  instant_bookable + neighbourhood

```

	Df	Sum of Sq	RSS	AIC
- cancellation_policy	1	0.062	367.75	-3568.1
- host_since	1	0.072	367.76	-3568.0
- host_has_profile_pic	1	0.155	367.84	-3567.5
- host_identity_verified	1	0.346	368.03	-3566.5
<none>			367.69	-3566.4
- property_type	1	0.578	368.27	-3565.1
- number_of_reviews	1	0.602	368.29	-3565.0
- host_response_rate	1	0.916	368.60	-3563.2
- bed_type	1	0.937	368.63	-3563.1
- neighbourhood	1	1.003	368.69	-3562.7
- review_scores_rating	1	1.682	369.37	-3558.9
- amenities	1	1.865	369.55	-3557.9
- instant_bookable	1	2.061	369.75	-3556.8
- cleaning_fee	1	2.136	369.82	-3556.4

```

- bathrooms          1      5.429 373.12 -3537.9
- longitude           1      8.408 376.10 -3521.4
- latitude            1     13.115 380.80 -3495.5
- accommodates        1     39.809 407.50 -3354.6
- room_type           1    119.998 487.69 -2980.9

```

Step: AIC=-3568.06

```

log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + host_has_profile_pic +
  host_identity_verified + host_since + instant_bookable +
  neighbourhood

```

	Df	Sum of Sq	RSS	AIC
- host_since	1	0.073	367.82	-3569.6
- host_has_profile_pic	1	0.163	367.91	-3569.1
- host_identity_verified	1	0.346	368.10	-3568.1
<none>			367.75	-3568.1
- property_type	1	0.593	368.34	-3566.7
- number_of_reviews	1	0.643	368.39	-3566.4
- host_response_rate	1	0.919	368.67	-3564.9
- bed_type	1	0.932	368.68	-3564.8
- neighbourhood	1	1.005	368.75	-3564.4
- review_scores_rating	1	1.756	369.51	-3560.1
- amenities	1	1.881	369.63	-3559.4
- instant_bookable	1	2.075	369.83	-3558.4
- cleaning_fee	1	2.630	370.38	-3555.2
- bathrooms	1	5.401	373.15	-3539.7
- longitude	1	8.384	376.13	-3523.2
- latitude	1	13.068	380.82	-3497.4
- accommodates	1	39.810	407.56	-3356.3
- room_type	1	120.521	488.27	-2980.4

Step: AIC=-3569.64

```

log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + host_has_profile_pic +
  host_identity_verified + instant_bookable + neighbourhood

```

	Df	Sum of Sq	RSS	AIC
- host_has_profile_pic	1	0.168	367.99	-3570.7

```

- host_identity_verified 1      0.319 368.14 -3569.8
<none>                      367.82 -3569.6
- property_type          1      0.603 368.43 -3568.2
- number_of_reviews      1      0.630 368.45 -3568.1
- host_response_rate     1      0.904 368.73 -3566.5
- bed_type               1      0.930 368.75 -3566.4
- neighbourhood          1      1.018 368.84 -3565.9
- review_scores_rating   1      1.740 369.56 -3561.8
- amenities              1      1.886 369.71 -3561.0
- instant_bookable       1      2.085 369.91 -3559.9
- cleaning_fee           1      2.633 370.46 -3556.8
- bathrooms              1      5.416 373.24 -3541.2
- longitude              1      8.340 376.16 -3525.0
- latitude               1     13.072 380.90 -3499.0
- accommodates           1     39.805 407.63 -3357.9
- room_type              1    120.614 488.44 -2981.7

```

Step: AIC=-3570.69

```

log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + host_identity_verified +
  instant_bookable + neighbourhood

```

	Df	Sum of Sq	RSS	AIC
- host_identity_verified	1	0.344	368.34	-3570.7
<none>			367.99	-3570.7
- property_type	1	0.597	368.59	-3569.3
- number_of_reviews	1	0.636	368.63	-3569.1
- host_response_rate	1	0.918	368.91	-3567.5
- bed_type	1	0.936	368.93	-3567.4
- neighbourhood	1	1.024	369.02	-3566.9
- review_scores_rating	1	1.756	369.75	-3562.8
- amenities	1	1.895	369.89	-3562.0
- instant_bookable	1	2.074	370.07	-3561.0
- cleaning_fee	1	2.671	370.66	-3557.6
- bathrooms	1	5.489	373.48	-3541.9
- longitude	1	8.343	376.33	-3526.1
- latitude	1	13.143	381.14	-3499.7
- accommodates	1	39.658	407.65	-3359.8
- room_type	1	120.765	488.76	-2982.4

Step: AIC=-3570.75

```
log_price ~ accommodates + bathrooms + cleaning_fee + host_response_rate +
  latitude + longitude + number_of_reviews + review_scores_rating +
  property_type + room_type + amenities + bed_type + instant_bookable +
  neighbourhood
```

	Df	Sum of Sq	RSS	AIC
<none>			368.34	-3570.7
- property_type	1	0.659	368.99	-3569.0
- number_of_reviews	1	0.784	369.12	-3568.3
- bed_type	1	0.872	369.21	-3567.8
- host_response_rate	1	0.997	369.33	-3567.1
- neighbourhood	1	1.024	369.36	-3567.0
- review_scores_rating	1	1.676	370.01	-3563.3
- instant_bookable	1	1.879	370.22	-3562.2
- amenities	1	1.890	370.23	-3562.1
- cleaning_fee	1	2.829	371.16	-3556.8
- bathrooms	1	5.501	373.84	-3541.9
- longitude	1	8.228	376.56	-3526.8
- latitude	1	13.575	381.91	-3497.5
- accommodates	1	39.493	407.83	-3360.9
- room_type	1	120.848	489.18	-2982.6

[Hide](#)

```
step_boston_mlr
```


Call:

```
lm(formula = log_price ~ accommodates + bathrooms + cleaning_fee +
    host_response_rate + latitude + longitude + number_of_reviews +
    review_scores_rating + property_type + room_type + amenities +
    bed_type + instant_bookable + neighbourhood, data = Training)
```

Coefficients:

(Intercept)	accommodates	bathrooms	cleaning_feeTRUE
1.201e+00	8.198e-02	1.230e-01	-9.052e-02
host_response_rate	latitude	longitude	number_of_reviews
-2.029e-01	3.824e+00	2.225e+00	-4.621e-04
review_scores_rating	property_type	room_type	amenities
4.012e-03	2.448e-03	-6.111e-01	1.780e-06
bed_type	instant_bookable	neighbourhood	
4.922e-02	-6.602e-02	1.312e-04	

Hide

```
summary(step_boston_mlr)
```

Call:

```
lm(formula = log_price ~ accommodates + bathrooms + cleaning_fee +
    host_response_rate + latitude + longitude + number_of_reviews +
    review_scores_rating + property_type + room_type + amenities +
    bed_type + instant_bookable + neighbourhood, data = Training)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.73087	-0.26077	-0.01602	0.26359	2.44244

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.201e+00	3.506e+01	0.034	0.97269
accommodates	8.198e-02	5.509e-03	14.880	< 2e-16 ***
bathrooms	1.230e-01	2.216e-02	5.553	3.16e-08 ***
cleaning_feeTRUE	-9.052e-02	2.273e-02	-3.982	7.06e-05 ***
host_response_rate	-2.029e-01	8.582e-02	-2.364	0.01817 *
latitude	3.824e+00	4.384e-01	8.724	< 2e-16 ***
longitude	2.225e+00	3.276e-01	6.792	1.44e-11 ***
number_of_reviews	-4.621e-04	2.204e-04	-2.096	0.03619 *
review_scores_rating	4.012e-03	1.309e-03	3.065	0.00220 **
property_type	2.448e-03	1.274e-03	1.922	0.05474 .
room_type	-6.111e-01	2.348e-02	-26.029	< 2e-16 ***
amenities	1.780e-06	5.470e-07	3.255	0.00115 **
bed_type	4.922e-02	2.226e-02	2.211	0.02716 *
instant_bookable	-6.602e-02	2.034e-02	-3.246	0.00119 **
neighbourhood	1.312e-04	5.476e-05	2.396	0.01667 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4223 on 2065 degrees of freedom

Multiple R-squared: 0.5994, Adjusted R-squared: 0.5967

F-statistic: 220.7 on 14 and 2065 DF, p-value: < 2.2e-16

Hide

```
boston_mlr_pred <- predict(boston_mlr, Validation)
accuracy(boston_mlr_pred, Validation$log_price)
```



```
'data.frame':  2080 obs. of  24 variables:
 $ log_price          : num  4.6 4.68 4.83 4.09 4.96 ...
 $ accommodates       : int   2 2 6 2 2 1 4 3 2 5 ...
 $ bathrooms          : num   2 1 1 1 1 1 1 1 1 1 ...
 $ latitude           : num  42.3 42.3 42.3 42.3 42.4 ...
 $ longitude          : num  -71 -71.1 -71.1 -71.1 -71.1 ...
 $ review_scores_rating : num   88 96 100 96 80 100 99 100 89 96 ...
 $ bedrooms           : int   1 1 2 1 1 1 2 2 1 2 ...
 $ beds              : num   1 1 4 1 1 1 2 2 1 2 ...
 $ id                 : int  14648556 4680055 4274462 2278299 16253186 14916417 8442997 2259813 7575345 982304
 ...
 $ property_type      : Factor w/ 35 levels "Apartment","Bed & Breakfast",...: 30 12 1 18 1 1 18 12 1 1 ...
 $ room_type          : Factor w/ 3 levels "Entire home/apt",...: 2 2 1 2 1 2 2 1 1 1 ...
 $ amenities          : Factor w/ 67122 levels "{ \"Air conditioning\", \"Carbon monoxide detector\", \"First aid
kit\", \"Lock on bedroom door\", \"translation mis\" | __truncated__,...: 13512 58394 48034 64902 50261 3940 17253 214
97 30934 20451 ...
 $ bed_type           : Factor w/ 5 levels "Airbed","Couch",...: 5 5 5 5 5 5 5 5 5 5 ...
 $ cancellation_policy : Factor w/ 5 levels "flexible","moderate",...: 3 3 3 1 1 1 1 1 4 1 ...
 $ cleaning_fee       : logi  TRUE TRUE TRUE FALSE FALSE FALSE ...
 $ description        : Factor w/ 73474 levels "          Cozy & clean on the corner of 5th Street and 2nd Ave
nue. 1 room (Queen bed) in 3 bedroom apartment (2\" | __truncated__,...: 62216 66393 6863 61398 41717 2057 49914 119
18 46947 8233 ...
 $ host_has_profile_pic : Factor w/ 3 levels "", "f", "t": 3 3 3 3 3 3 3 3 3 3 ...
 $ host_identity_verified: Factor w/ 3 levels "", "f", "t": 3 3 3 2 2 2 3 3 3 3 ...
 $ host_response_rate   : num   1 1 1 1 1 1 1 1 0.88 1 ...
 $ host_since          : Factor w/ 3088 levels "", "1/1/11", "1/1/12",...: 2442 2078 141 785 1979 2563 468 3004 22
80 1723 ...
 $ instant_bookable     : Factor w/ 2 levels "f", "t": 1 2 1 1 2 1 1 1 1 1 ...
 $ name                : Factor w/ 73350 levels " 1 Bed Apt in Utopic Williamsburg ",...: 20723 51268 493 9707 1
9414 2583 69112 14037 334 43075 ...
 $ neighbourhood        : Factor w/ 620 levels "", "16th Street Heights",...: 500 500 500 281 44 359 591 100 503 8
...
 $ number_of_reviews     : int   12 40 5 0 2 1 61 5 9 0 ...
```

Hide

```
# Min & Max of Predictor values
names(bos)
```

```
[1] "log_price"          "accommodates"      "bathrooms"        "cleaning_fee"
[5] "host_response_rate" "latitude"          "longitude"        "number_of_reviews"
[9] "review_scores_rating" "property_type"     "room_type"        "amenities"
[13] "bed_type"          "cancellation_policy" "description"      "host_has_profile_pic"
[17] "host_identity_verified" "host_since"       "instant_bookable" "name"
[21] "neighbourhood"
```

Hide

```
accommodates <- runif(1, min(bostontrain$accommodates), max(bostontrain$accommodates))
bathrooms <- runif(1, min(bostontrain$bathrooms), max(bostontrain$bathrooms))
bedrooms<- runif(1, min(bostontrain$bedrooms), max(bostontrain$bedrooms))
beds<- runif(1, min(bostontrain$beds), max(bostontrain$beds))
log_price <- runif(1, min(bostontrain$log_price), max(bostontrain$log_price))
review_scores_rating <- runif(1, min(bostontrain$review_scores_rating), max(bostontrain$review_scores_rating))
latitude <- runif(1, min(bostontrain$latitude), max(bostontrain$latitude))
longitude <- runif(1, min(bostontrain$longitude), max(bostontrain$longitude))

names(bostontrain)
```

```
[1] "log_price"          "accommodates"      "bathrooms"        "latitude"
[5] "longitude"          "review_scores_rating" "bedrooms"         "beds"
[9] "id"                 "property_type"     "room_type"        "amenities"
[13] "bed_type"          "cancellation_policy" "cleaning_fee"     "description"
[17] "host_has_profile_pic" "host_identity_verified" "host_response_rate" "host_since"
[21] "instant_bookable"  "name"              "neighbourhood"    "number_of_reviews"
```

Hide

```
log_price
```

```
[1] 3.982864
```

Hide

```
accommodates
```

```
[1] 14.47223
```

[Hide](#)

```
bathrooms
```

```
[1] 1.777626
```

[Hide](#)

```
latitude
```

```
[1] 42.33326
```

[Hide](#)

```
longitude
```

```
[1] -70.9895
```

[Hide](#)

```
review_scores_rating
```

```
[1] 71.25128
```

[Hide](#)

```
bedrooms
```

```
[1] 0.1491482
```

[Hide](#)

```
beds
```

```
[1] 12.2405
```

Hide

```
# Creating rental_fee dataframe
colnames(bostontrain)
```

```
[1] "log_price"          "accommodates"      "bathrooms"         "latitude"
[5] "longitude"          "review_scores_rating" "bedrooms"          "beds"
[9] "id"                 "property_type"      "room_type"         "amenities"
[13] "bed_type"           "cancellation_policy" "cleaning_fee"       "description"
[17] "host_has_profile_pic" "host_identity_verified" "host_response_rate" "host_since"
[21] "instant_bookable"    "name"               "neighbourhood"      "number_of_reviews"
```

Hide

```
rental_fee <- data.frame(log_price=5.89,
                        accommodates=11.0,
                        bathrooms=1.5,
                        latitude=42.26,
                        longitude=-71.0,
                        review_scores_rating=26.0,
                        bedrooms=3.0,
                        beds=11.0)
```

```
train.norm <- bostontrain
valid.norm <- bostonvalid
rental.norm <- rental1
```

```
install.packages('caret')
```

```
Error in install.packages : Updating loaded packages
```

Hide

```
library(caret)
norm.values <- preProcess(bostontrain[, 2:8], method=c("center", "scale"))
train.norm[, 2:8] <- predict(norm.values, bostontrain[, 2:8])
valid.norm[, 2:8] <- predict(norm.values, bostonvalid[, 2:8])
rental.norm[, 2:8] <- predict(norm.values, rental1[, 2:8])
new.norm <- predict(norm.values, rental_fee)
# Use Knn Function to find nearest neighbors
install.packages("FNN")
```

```
Installing package into '/Users/josemartinez/Library/R/3.6/library'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/macosx/el-capitan/contrib/3.6/FNN_1.1.3.tgz'
Content type 'application/x-gzip' length 136106 bytes (132 KB)
=====
downloaded 132 KB
```

The downloaded binary packages are in
 /var/folders/6v/wsr694r57n9dfsdxfthdysbh0000gn/T//RtmplIEjsq/downloaded_packages

[Hide](#)

```
install.packages("caret")
```

```
Installing package into '/Users/josemartinez/Library/R/3.6/library'
(as 'lib' is unspecified)
trying URL 'https://cran.rstudio.com/bin/macosx/el-capitan/contrib/3.6/caret_6.0-86.tgz'
Content type 'application/x-gzip' length 6252869 bytes (6.0 MB)
=====
downloaded 6.0 MB
```

The downloaded binary packages are in
 /var/folders/6v/wsr694r57n9dfsdxfthdysbh0000gn/T//RtmplIEjsq/downloaded_packages

[Hide](#)


```
library(FNN)
nn <- knn(train = train.norm[, 2:8], test = new.norm[, 2:8],
         cl=train.norm[, 15], k=9)
nn
```

```
[1] TRUE
attr(,"nn.index")
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9]
[1,]  936 2003 1221 1848 1006   52  424  565  170
attr(,"nn.dist")
      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]      [,8]      [,9]
[1,] 6.68265 8.727249 8.970678 9.055821 9.148725 9.170462 9.504666 9.809327 9.900007
Levels: TRUE
```

Hide

```
row.names(bostontrain)[attr(nn, "nn.index")]
```

```
[1] "936" "2003" "1221" "1848" "1006" "52" "424" "565" "170"
```

Hide

```
# Accuracy
accuracy.rental<- data.frame(k = seq(1, 10, 1), accuracy = rep(0, 10))
for(i in 1:10) {
  knn.pred <- knn(train.norm[, 2:8], valid.norm[, 2:8],
                  cl = train.norm[, 10], k = i)
  accuracy.rental[i, 2] <- confusionMatrix(knn.pred, valid.norm[, 10])$overall[1]
}
```



```
library(e1071)
# Part C
boston2 <- boston1 # copy boston 1 df to preserve original data
summary(boston2$log_price) # capture the range
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
2.833	4.382	4.913	4.884	5.298	7.244

Hide

```
summary(boston2$nightly_price) # capture the nightly price and compare to log
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
17.0	80.0	136.0	165.6	200.0	1400.0

Hide

```
# create bins using the cut function to create 4 price categories
boston2$log_price_rating <- cut(boston2$log_price,
                               breaks=c(2.833, 4.382, 5.298, 7.244, Inf),
                               labels=c("Student Budget", "Below Average", "Above Average", "Pricey Dig"))
summary(boston2$log_price_rating)
```

Student Budget	Below Average	Above Average	Pricey Dig
827	1739	901	1

Hide

```
# Part D
boston2$log_price <- factor(boston2$log_price)
boston2$property_type <- factor(boston2$property_type)
boston2$cancellation_policy <- factor(boston2$cancellation_policy)
boston2$bed_type <- factor(boston2$bed_type)
boston2$cleaning_fee <- factor(boston2$cleaning_fee)
boston2$log_price_rating <- factor(boston2$log_price_rating)
# create training and validation sets
selected.var <- c(2, 8, 9, 10, 26)
train.index <- sample(c(1:dim(boston2)[1]), dim(boston2)[1]*.60)
boston2_train <- boston2[train.index, selected.var]
boston2_val <- boston2[-train.index, selected.var]
# run naive bayes
boston2.nb <- naiveBayes(log_price_rating ~., data = boston2_train)
boston2.nb
```

Naive Bayes Classifier for Discrete Predictors

Call:

naiveBayes.default(x = X, y = Y, laplace = laplace)

A-priori probabilities:

Y

Student Budget	Below Average	Above Average	Pricey Dig
0.2471153846	0.4985576923	0.2538461538	0.0004807692

Conditional probabilities:

	log_price							
Y	2.833213344	2.995732274	3.091042453	3.17805383	3.218875825	3.295836866	3.33220451	
Student Budget	0.0000000000	0.0019455253	0.0019455253	0.0019455253	0.0058365759	0.0038910506	0.0058365759	
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
	log_price							
Y	3.36729583	3.401197382	3.433987204	3.465735903	3.496507561	3.526360525	3.555348061	
Student Budget	0.0097276265	0.0214007782	0.0077821012	0.0077821012	0.0097276265	0.0019455253	0.0252918288	
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
	log_price							
Y	3.583518938	3.610917913	3.63758616	3.663561646	3.688879454	3.713572067	3.737669618	
Student Budget	0.0000000000	0.0058365759	0.0097276265	0.0175097276	0.0311284047	0.0097276265	0.0038910506	
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
	log_price							
Y	3.784189634	3.80666249	3.828641396	3.850147602	3.871201011	3.891820298	3.912023005	
Student Budget	0.0077821012	0.0272373541	0.0077821012	0.0097276265	0.0038910506	0.0136186770	0.0836575875	
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
	log_price							
Y	3.931825633	3.951243719	3.970291914	3.988984047	4.007333185	4.025351691	4.043051268	
Student Budget	0.0077821012	0.0038910506	0.0077821012	0.0058365759	0.0603112840	0.0077821012	0.0058365759	
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	
	log_price							
Y	4.060443011	4.077537444	4.094344562	4.110873864	4.127134385	4.143134726	4.158883083	
Student Budget	0.0116731518	0.0175097276	0.0778210117	0.0058365759	0.0116731518	0.0175097276	0.0077821012	

	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.17438727	4.189654742	4.204692619	4.219507705	4.234106505	4.248495242	4.262679877
	Student Budget	0.0914396887	0.0019455253	0.0194552529	0.0097276265	0.0505836576	0.0642023346	0.0097276265
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.276666119	4.290459441	4.304065093	4.317488114	4.330733334	4.343805422	4.356708827
	Student Budget	0.0097276265	0.0019455253	0.0116731518	0.0953307393	0.0058365759	0.0019455253	0.0077821012
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.369447852	4.382026635	4.394449155	4.406719247	4.418840608	4.430816799	4.442651256
	Student Budget	0.0330739300	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0327868852	0.0019286403	0.0028929605	0.0019286403	0.0038572806	0.0241080039
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.454347296	4.477336814	4.48863637	4.49980967	4.510859507	4.521788577	4.532599493
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0038572806	0.0096432015	0.0154291225	0.0270009643	0.0009643202	0.0009643202	0.0028929605
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.543294782	4.553876892	4.564348191	4.574710979	4.584967479	4.59511985	4.605170186
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0009643202	0.0260366442	0.0009643202	0.0048216008	0.0067502411	0.0453230473	0.0520732883
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.624972813	4.634728988	4.644390899	4.65396035	4.663439094	4.672828834	4.682131227
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0009643202	0.0028929605	0.0038572806	0.0077145612	0.0009643202	0.0009643202	0.0067502411
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.691347882	4.700480366	4.709530201	4.718498871	4.727387819	4.736198448	4.744932128
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0086788814	0.0231436837	0.0019286403	0.0019286403	0.0009643202	0.0028929605	0.0135004822
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.753590191	4.762173935	4.770684624	4.779123493	4.787491743	4.795790546	4.804021045
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0019286403	0.0009643202	0.0019286403	0.0096432015	0.0270009643	0.0009643202	0.0028929605

	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.812184355	4.820281566	4.828313737	4.836281907	4.844187086	4.852030264	4.859812404
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0028929605	0.0405014465	0.0019286403	0.0028929605	0.0028929605	0.0395371263
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.86753445	4.882801923	4.890349128	4.8978398	4.905274778	4.912654886	4.919980926
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0202507232	0.0019286403	0.0019286403	0.0038572806	0.0192864031	0.0028929605	0.0019286403
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.927253685	4.934473933	4.941642423	4.94875989	4.955827058	4.96284463	4.9698133
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0019286403	0.0144648023	0.0163934426	0.0000000000	0.0009643202	0.0009643202	0.0057859209
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		4.976733742	4.983606622	4.990432587	4.997212274	5.003946306	5.010635294	5.017279837
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0096432015	0.0009643202	0.0086788814	0.0038572806	0.0241080039	0.0597878496	0.0009643202
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		5.023880521	5.030437921	5.036952602	5.043425117	5.049856007	5.056245805	5.062595033
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0038572806	0.0009643202	0.0019286403	0.0067502411	0.0009643202	0.0009643202	0.0048216008
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		5.068904202	5.075173815	5.081404365	5.087596335	5.093750201	5.099866428	5.105945474
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0144648023	0.0183220829	0.0000000000	0.0009643202	0.0000000000	0.0009643202	0.0077145612
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		5.111987788	5.117993812	5.123963979	5.129898715	5.135798437	5.141663557	5.147494477
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0019286403	0.0028929605	0.0009643202	0.0135004822	0.0144648023	0.0009643202	0.0028929605
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price							
Y		5.153291594	5.159055299	5.164785974	5.170483995	5.176149733	5.18178355	5.187385806
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0009643202	0.0048216008	0.0540019286	0.0067502411	0.0028929605	0.0009643202	0.0173577628
	Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000

	log_price						
Y	5.192956851	5.198497031	5.204006687	5.209486153	5.214935758	5.220355825	5.225746674
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0163934426	0.0028929605	0.0000000000	0.0009643202	0.0038572806	0.0163934426	0.0038572806
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price						
Y	5.231108617	5.236441963	5.241747015	5.247024072	5.252273428	5.257495372	5.262690189
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0019286403	0.0038572806	0.0173577628	0.0106075217	0.0019286403	0.0009643202	0.0019286403
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	log_price						
Y	5.272999559	5.278114659	5.283203729	5.288267031	5.293304825	5.298317367	5.303304908
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0183220829	0.0019286403	0.0000000000	0.0019286403	0.0279652845	0.0000000000	0.0000000000
Above Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.1212121212	0.0056818182
	log_price						
Y	5.313205979	5.318119994	5.323009979	5.332718793	5.33753808	5.342334252	5.347107531
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0000000000	0.0018939394	0.0056818182	0.0037878788	0.0000000000	0.0113636364	0.0170454545
	log_price						
Y	5.356586275	5.361292166	5.370638028	5.375278408	5.384495063	5.38907173	5.393627546
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0000000000	0.0018939394	0.0075757576	0.0075757576	0.0000000000	0.0113636364	0.0189393939
	log_price						
Y	5.407171771	5.416100402	5.429345629	5.433722004	5.438079309	5.455321115	5.459585514
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0037878788	0.0492424242	0.0000000000	0.0151515152	0.0075757576	0.0018939394	0.0075757576
	log_price						
Y	5.463831805	5.468060141	5.476463552	5.480638923	5.501258211	5.505331536	5.509388337
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0000000000	0.0037878788	0.0037878788	0.0113636364	0.0075757576	0.0037878788	0.0056818182
	log_price						
Y	5.517452896	5.521460918	5.525452939	5.541263545	5.549076085	5.556828062	5.560681631
Student Budget	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000	0.000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0265151515	0.0776515152	0.0018939394	0.0056818182	0.0018939394	0.0208333333	0.0113636364
	log_price						

Y	5.564520407	5.575949103	5.579729826	5.587248658	5.594711138	5.598421959	5.616771098
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0018939394	0.0018939394	0.0037878788	0.0037878788	0.0018939394	0.0037878788	0.0359848485
	log_price						
Y	5.620400866	5.631211782	5.634789603	5.638354669	5.648974238	5.65248918	5.655991811
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0000000000	0.0056818182	0.0056818182	0.0018939394	0.0018939394	0.0056818182	0.0018939394
	log_price						
Y	5.659482216	5.66296048	5.666426688	5.669880923	5.673323267	5.676753802	5.680172609
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0000000000	0.0018939394	0.0037878788	0.0037878788	0.0018939394	0.0037878788	0.0000000000
	log_price						
Y	5.683579767	5.686975356	5.697093487	5.700443573	5.703782475	5.720311777	5.723585102
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0018939394	0.0151515152	0.0056818182	0.0208333333	0.0435606061	0.0018939394	0.0018939394
	log_price						
Y	5.726847748	5.739792912	5.752572639	5.758901774	5.765191103	5.768320996	5.780743516
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0018939394	0.0018939394	0.0056818182	0.0018939394	0.0018939394	0.0056818182	0.0037878788
	log_price						
Y	5.783825182	5.789960171	5.793013608	5.796057751	5.799092654	5.802118375	5.814130532
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0132575758	0.0018939394	0.0000000000	0.0037878788	0.0000000000	0.0000000000	0.0037878788
	log_price						
Y	5.82008293	5.826000107	5.831882477	5.834810737	5.843544417	5.84932478	5.855071922
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0018939394	0.0018939394	0.0037878788	0.0018939394	0.0000000000	0.0000000000	0.0094696970
	log_price						
Y	5.857933154	5.863631176	5.877735782	5.883322388	5.886104031	5.888877958	5.894402834
Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
Above Average	0.0435606061	0.0018939394	0.0018939394	0.0018939394	0.0037878788	0.0000000000	0.0000000000
	log_price						
Y	5.908082938	5.910796644	5.926926026	5.937536205	5.940171253	5.948034989	5.963579344

	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0018939394	0.0018939394	0.0075757576	0.0018939394	0.0018939394	0.0000000000	0.0037878788
	log_price							
Y		5.966146739	5.96870756	5.978885765	5.988961417	5.991464547	6.033086222	6.052089169
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0000000000	0.0018939394	0.0075757576	0.0151515152	0.0208333333	0.0037878788	0.0056818182
	log_price							
Y		6.061456919	6.084499413	6.109247583	6.11368218	6.163314804	6.171700597	6.173786104
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0075757576	0.0037878788	0.0151515152	0.0018939394	0.0018939394	0.0018939394	0.0018939394
	log_price							
Y		6.184148891	6.204557763	6.212606096	6.214608098	6.222576268	6.224558429	6.226536669
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0018939394	0.0018939394	0.0018939394	0.0170454545	0.0000000000	0.0018939394	0.0018939394
	log_price							
Y		6.242223265	6.251903883	6.261491684	6.263398263	6.284134161	6.308098442	6.309918278
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0037878788	0.0018939394	0.0000000000	0.0018939394	0.0018939394	0.0018939394	0.0018939394
	log_price							
Y		6.324358962	6.326149473	6.342121419	6.34738921	6.354370041	6.388561406	6.391917113
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0000000000	0.0000000000	0.0018939394	0.0094696970	0.0018939394	0.0000000000	0.0000000000
	log_price							
Y		6.395261598	6.396929655	6.401917197	6.43775165	6.456769656	6.476972363	6.522092798
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0000000000	0.0018939394	0.0000000000	0.0037878788	0.0018939394	0.0056818182	0.0018939394
	log_price							
Y		6.549650742	6.551080335	6.586171655	6.620073207	6.64509097	6.65929392	6.665683718
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Below Average	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000
	Above Average	0.0075757576	0.0056818182	0.0000000000	0.0056818182	0.0000000000	0.0018939394	0.0000000000
	log_price							
Y		6.684611728	6.725033642	6.770789424	6.802394763	6.817830571	6.902742737	6.906754779
	Student Budget	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000	0.0000000000

```

Below Average  0.0000000000 0.0000000000 0.0000000000 0.0000000000 0.0000000000 0.0000000000 0.0000000000
Above Average  0.0113636364 0.0000000000 0.0000000000 0.0018939394 0.0018939394 0.0018939394 0.0037878788
log_price
Y
7.13089883 7.150701458 7.244227516
Student Budget 0.0000000000 0.0000000000 0.0000000000
Below Average  0.0000000000 0.0000000000 0.0000000000
Above Average  0.0000000000 0.0018939394 0.0000000000
[ reached getOption("max.print") -- omitted 1 row ]

```

```

bed_type
Y
Airbed Couch Futon Pull-out Sofa Real Bed
Student Budget 0.027237354 0.003891051 0.019455253 0.015564202 0.933852140
Below Average  0.002892960 0.001928640 0.001928640 0.001928640 0.991321119
Above Average  0.003787879 0.001893939 0.000000000 0.000000000 0.994318182
Pricey Dig 0.000000000 0.000000000 0.000000000 0.000000000 1.000000000

cancellation_policy
Y
flexible moderate strict super_strict_30 super_strict_60
Student Budget 0.350194553 0.239299611 0.410505837 0.000000000 0.000000000
Below Average  0.217936355 0.238187078 0.538090646 0.005785921 0.000000000
Above Average  0.176136364 0.236742424 0.556818182 0.028409091 0.001893939
Pricey Dig 0.000000000 1.000000000 0.000000000 0.000000000 0.000000000

```

```

cleaning_fee
Y
FALSE TRUE
Student Budget 0.3093385 0.6906615
Below Average  0.2121504 0.7878496
Above Average  0.2026515 0.7973485
Pricey Dig 1.0000000 0.0000000

```

[Hide](#)

```

# create a prop table
prop.table(table(boston2_train$cancellation_policy, boston2_train$log_price_rating), margin = 2)

```

	Student Budget	Below Average	Above Average	Pricey Dig
flexible	0.350194553	0.217936355	0.176136364	0.000000000
moderate	0.239299611	0.238187078	0.236742424	1.000000000
strict	0.410505837	0.538090646	0.556818182	0.000000000
super_strict_30	0.000000000	0.005785921	0.028409091	0.000000000
super_strict_60	0.000000000	0.000000000	0.001893939	0.000000000

Hide

```

pred.prob <- predict(boston2.nb, newdata = boston2_val, type = "raw")
pred.class <- predict(boston2.nb, newdata = boston2_val)
boston2_df <- data.frame(actual = boston2_val$log_price_rating, predicted = pred.class, pred.prob)
# dummy predictor for apartment type rental with a real bed, above average rating, and flexible cancellation
boston2_df[boston2_val$bed_type == "Real Bed" &
  boston2_val$cancellation_policy == "flexible" &
  boston2_val$cleaning_fee == "TRUE" &
  boston2_val$log_price_rating == "Student Budget", ]

```

	actual <ctr>	predicted <ctr>	Student.Budget <dbl>	Below.Average <dbl>	Above.Average <dbl>	Pricey.Dig <dbl>
74	Student Budget	Student Budget	0.7302980	0.19023578	0.079466231	0.00000000107777005
126	Student Budget	Student Budget	0.9748818	0.01771727	0.007400946	0.00000000010037620
172	Student Budget	Student Budget	0.7830988	0.15299243	0.063908753	0.000000000086676994
178	Student Budget	Student Budget	0.8633542	0.09638386	0.040261943	0.000000000054605731
202	Student Budget	Student Budget	0.9748818	0.01771727	0.007400946	0.00000000010037620
205	Student Budget	Student Budget	0.9769703	0.01624414	0.006785584	0.00000000009203028
209	Student Budget	Student Budget	0.9214687	0.05539245	0.023138810	0.000000000031382281
326	Student Budget	Student Budget	0.8186103	0.12794419	0.053445480	0.000000000072486057
331	Student Budget	Student Budget	0.8903916	0.07731289	0.032295522	0.000000000043801179
383	Student Budget	Student Budget	0.9769703	0.01624414	0.006785584	0.00000000009203028

1-10 of 50 rows

[Previous](#) **1** [2](#) [3](#) [4](#) [5](#) [Next](#)[Hide](#)

```
# Assesing model
library(caret)
pred.class <- predict(boston2.nb, newdata = boston2_train)
confusionMatrix(pred.class, boston2_train$log_price_rating)
```

Confusion Matrix and Statistics

Prediction	Reference			
	Student Budget	Below Average	Above Average	Pricey Dig
Student Budget	512	6	3	0
Below Average	2	1030	16	0
Above Average	0	1	509	0
Pricey Dig	0	0	0	1

Overall Statistics

Accuracy : 0.9865
 95% CI : (0.9806, 0.991)
 No Information Rate : 0.4986
 P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.9784

Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: Student Budget	Class: Below Average	Class: Above Average	Class: Pricey Dig
Sensitivity	0.9961	0.9932	0.9640	1.0000000
Specificity	0.9943	0.9827	0.9994	1.0000000
Pos Pred Value	0.9827	0.9828	0.9980	1.0000000
Neg Pred Value	0.9987	0.9932	0.9879	1.0000000
Prevalence	0.2471	0.4986	0.2538	0.0004808
Detection Rate	0.2462	0.4952	0.2447	0.0004808
Detection Prevalence	0.2505	0.5038	0.2452	0.0004808
Balanced Accuracy	0.9952	0.9880	0.9817	1.0000000

[Hide](#)

```

pred.class <- predict(boston2.nb, newdata = boston2_val)
confusionMatrix(pred.class, boston2_val$log_price_rating)

```

Confusion Matrix and Statistics

Prediction	Reference			
	Student Budget	Below Average	Above Average	Pricey Dig
Student Budget	306	1	11	0
Below Average	7	701	47	0
Above Average	0	0	315	0
Pricey Dig	0	0	0	0

Overall Statistics

Accuracy : 0.9524
 95% CI : (0.9399, 0.963)
 No Information Rate : 0.5058
 P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.9223

Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: Student Budget	Class: Below Average	Class: Above Average	Class: Pricey Dig
Sensitivity	0.9776	0.9986	0.8445	NA
Specificity	0.9888	0.9213	1.0000	1
Pos Pred Value	0.9623	0.9285	1.0000	NA
Neg Pred Value	0.9935	0.9984	0.9459	NA
Prevalence	0.2255	0.5058	0.2687	0
Detection Rate	0.2205	0.5050	0.2269	0
Detection Prevalence	0.2291	0.5439	0.2269	0
Balanced Accuracy	0.9832	0.9599	0.9223	NA

Part III. Classification Tree

[Code](#)

```
[1] "The process of building a Classification Tree that predicts the outcome of the cancelation policy of an Airbnb rental listing involved several steps. After importing and filtering the data into a Boston-only dataframe, we converted listings from the dataset where the cancelation policy was either super_strict_60 or super_strict_30 in to just "strict". Since the assignment was to predict the outcome cancelation policy into only one of three buckets (flexible, moderate, or strict) these other cancelation policies needed to be removed. Luckily there were very few observations within the Boston dataset that had the super_strict policy. The next step was to replace NULL values in the dataset with NA, and then impute median values to replace the NA values. This ensured that our calculations would function as intended, since the packages we were utilizing cannot handle NULL or NA values in their computations. Next, we renamed the row names as the ID field of the Airbnb listing, and removed columns that we did not intend to use as a part of the classification tree. We chose the rows to remove by considering the value that they would off to the model, while also contemplating the computational strain that certain categorical variables would cause to our local machines. While the rpart() and rpartplot() packages can handle categorical variables, ones with many different values across the observations in the dataset are computationally cumbersome and the resulting number of splits in the tree can be unwieldy. We ultimately chose to include 12 input variables, most of which were numeric, and a few more manageable categorical values like "Bed_Type" where there were only a few possible values. After that, we confirmed the 12 remaining variables were formatted in the datatype we desired and then partitioned out data into training and testing datasets.\n\nOnce the training data was ready, we utilized the rpart() package with a complexity parameter of 0 to build our initial tree. By setting the CP = 0, we were ensuring that rpart() would build a massive tree that would surely overfit the training data. While this may seem like a wasted step since we knew we would not utilize the resulting tree, by doing this we were able to run the printcp() function and find the ideal number of splits for our tree where we would minimize the cross-validation error (denoted as xerror in the console). We then plotted the CP = 0 tree to confirm our suspicions that the tree was in fact too large. After plotting the tree and finding the CP value that corresponded to the lowest xerror, we reran our rpart() function with the xerror minimizing CP value, which resulted in a new classification tree with an ideal number of splits. For our dataset, the CP value that corresponded to the lowest xerror was 0.00501. \n"
```

[Hide](#)

```
library(reshape2)
```


Attaching package: 'reshape2'

The following objects are masked from 'package:reshape':

colsplit, melt, recast

The following object is masked from 'package:tidyr':

smiths

[Hide](#)

```
library(caret)
library(rpart)
library(rpart.plot)

head(df)
```

	id <int>	log_price <dbl>	property_type <fctr>	room_type <fctr>	
1	6901257	5.010635	Apartment	Entire home/apt	
2	6304928	5.129899	Apartment	Entire home/apt	
3	7919400	4.976734	Apartment	Entire home/apt	
4	13418779	6.620073	House	Entire home/apt	
5	3808709	4.744932	Apartment	Entire home/apt	
6	12422935	4.442651	Apartment	Private room	
6 rows 1-5 of 29 columns					

[Hide](#)

```

set.seed(200)
boston3 <- filter(df, city=="Boston")
# Convert 'super_strict_60' & 'super_strict_30' to just 'strict'
boston3 <- data.frame(lapply(boston3, function(x) {gsub("super_", "", x) })))
boston3 <- data.frame(lapply(boston3, function(x) {gsub("_30", "", x) })))
boston3 <- data.frame(lapply(boston3, function(x) {gsub("_60", "", x) })))
# check for 'NA' values
sum(is.na(boston3))

```

```
[1] 659
```

Hide

```
colSums(is.na(boston3))
```

id	log_price	property_type	room_type
0	0	0	0
amenities	accommodates	bathrooms	bed_type
0	0	6	0
cancellation_policy	cleaning_fee	city	description
0	0	0	0
first_review	host_has_profile_pic	host_identity_verified	host_response_rate
0	0	0	0
host_since	instant_bookable	last_review	latitude
0	0	0	0
longitude	name	neighbourhood	number_of_reviews
0	0	0	0
review_scores_rating	thumbnail_url	zipcode	bedrooms
648	0	0	3
beds			
2			

Hide

```
# Convert null values to 'NA'
boston3[boston3== ""] <-NA
# Impute median values when 'NA' & convert host_response_rate to numeric
boston3$review_scores_rating <- as.numeric(boston3$review_scores_rating)
boston3$review_scores_rating[is.na(boston3$review_scores_rating)] <- median(boston3$review_scores_rating, na.rm = TRUE)
boston3$host_response_rate <- as.numeric(sub("%","",boston3$host_response_rate))/100
boston3$host_response_rate[is.na(boston3$host_response_rate)] <- median(boston3$host_response_rate, na.rm=TRUE)
boston3$beds <- as.numeric(boston3$beds)
boston3$beds[is.na(boston3$beds)] <- median(boston3$beds, na.rm = T)
boston3$bathrooms <- as.numeric(boston3$bathrooms)
boston3$bathrooms[is.na(boston3$bathrooms)] <- median(boston3$bathrooms, na.rm=T)
boston3$bedrooms <-as.numeric(boston3$bedrooms)
boston3$bedrooms[is.na(boston3$bedrooms)] <- median(boston3$bedrooms, na.rm=T)

colSums(is.na(boston3))
```

id	log_price	property_type	room_type
0	0	0	0
amenities	accommodates	bathrooms	bed_type
0	0	0	0
cancellation_policy	cleaning_fee	city	description
0	0	0	0
first_review	host_has_profile_pic	host_identity_verified	host_response_rate
621	0	0	0
host_since	instant_bookable	last_review	latitude
0	0	621	0
longitude	name	neighbourhood	number_of_reviews
0	0	0	0
review_scores_rating	thumbnail_url	zipcode	bedrooms
0	134	26	0
beds			
0			

[Hide](#)

```
sum(is.na(boston3))
```

```
[1] 1402
```

Hide

```
# rename rows as id field
rownames(boston3) <- boston3[,1]
# remove fields we do not intent to use
names(boston3)
```

```
[1] "id"                "log_price"          "property_type"      "room_type"
[5] "amenities"         "accommodates"       "bathrooms"          "bed_type"
[9] "cancellation_policy" "cleaning_fee"       "city"               "description"
[13] "first_review"      "host_has_profile_pic" "host_identity_verified" "host_response_rate"
[17] "host_since"        "instant_bookable"   "last_review"        "latitude"
[21] "longitude"         "name"              "neighbourhood"      "number_of_reviews"
[25] "review_scores_rating" "thumbnail_url"      "zipcode"            "bedrooms"
[29] "beds"
```

Hide

```
boston3 <- boston3[-c(1, 3, 5, 11, 12, 13, 14, 15, 16, 17, 19, 20, 21, 22, 23, 26, 27)]
# confirm numeric variables are in fact numeric, convert if not
class(boston3$log_price)
```

```
[1] "factor"
```

Hide

```
boston3$log_price <- as.numeric((boston3$log_price))
class(boston3$room_type) #factor
```

```
[1] "factor"
```

Hide

```
class(boston3$accommodates)
```

```
[1] "factor"
```

[Hide](#)

```
boston3$accommodates <- as.numeric(boston3$accommodates)
class(boston3$bathrooms)
```

```
[1] "numeric"
```

[Hide](#)

```
class(boston3$bed_type) #factor
```

```
[1] "factor"
```

[Hide](#)

```
class(boston3$cancellation_policy) #factor
```

```
[1] "factor"
```

[Hide](#)

```
class(boston3$cleaning_fee) #factor
```

```
[1] "factor"
```

[Hide](#)

```
class(boston3$instant_bookable) #factor
```

```
[1] "factor"
```

[Hide](#)

```
class(boston3$number_of_reviews)
```

```
[1] "factor"
```

[Hide](#)

```
boston3$number_of_reviews <- as.numeric(boston3$number_of_reviews)
class(boston3$review_scores_rating)
```

```
[1] "numeric"
```

[Hide](#)

```
class(boston3$bedrooms)
```

```
[1] "numeric"
```

[Hide](#)

```
class(boston3$beds)
```

```
[1] "numeric"
```

[Hide](#)

```
# Create data partition of dataset
train.index <- createDataPartition(boston3$cancellation_policy,
                                   p = 0.60, #percentage split, enter desired portion for training data (60/40 split)
                                   list = FALSE, #tells it that we do not want it to come out as a list
                                   times = 1)

bos3_train <- boston3[train.index ,]
bos3_valid <- boston3[-train.index ,]

names(bos3_train)
```

```
[1] "log_price"      "room_type"      "accommodates"   "bathrooms"
[5] "bed_type"       "cancellation_policy" "cleaning_fee"    "instant_bookable"
[9] "number_of_reviews" "review_scores_rating" "bedrooms"       "beds"
```

[Hide](#)

```
names(bos3_valid)
```

```
[1] "log_price"      "room_type"      "accommodates"   "bathrooms"
[5] "bed_type"       "cancellation_policy" "cleaning_fee"    "instant_bookable"
[9] "number_of_reviews" "review_scores_rating" "bedrooms"       "beds"
```

[Hide](#)

```
# CREATING CLASSIFICATION TREES WITH VARIOUS COMPLEXITY PARAMETERS
options(scipen = 999)
# rpart with cp = 0 --- creates unpruned very large tree
tree_bos3_cp0 <- rpart(cancellation_policy ~ .,
                      data=bos3_train,
                      method = "class",
                      xval = 5,
                      cp = 0)
printcp(tree_bos3_cp0)
```

Classification tree:

```
rpart(formula = cancellation_policy ~ ., data = bos3_train, method = "class",
      xval = 5, cp = 0)
```

Variables actually used in tree construction:

[1] accommodates	bathrooms	bed_type	bedrooms	cleaning_fee
[6] instant_bookable	log_price	number_of_reviews	review_scores_rating	room_type


Root node error: 998/2083 = 0.47912

n= 2083

	CP	nsplit	rel error	xerror	xstd
1	0.1272545	0	1.00000	1.00000	0.022846
2	0.0115230	1	0.87275	0.87275	0.022557
3	0.0100200	3	0.84970	0.86874	0.022542
4	0.0046760	4	0.83968	0.85872	0.022504
5	0.0035070	7	0.82565	0.87375	0.022561
6	0.0027555	9	0.81864	0.88878	0.022613
7	0.0025050	14	0.80461	0.89579	0.022635
8	0.0020040	29	0.75150	0.90581	0.022666
9	0.0017307	44	0.71944	0.91283	0.022685
10	0.0015030	60	0.68838	0.90982	0.022677
11	0.0013360	66	0.67936	0.90281	0.022657
12	0.0010020	72	0.67134	0.92285	0.022712
13	0.0006680	89	0.65331	0.93587	0.022744
14	0.0005010	98	0.64729	0.94890	0.022771
15	0.0002505	104	0.64429	0.95892	0.022790
16	0.0002004	108	0.64329	0.96794	0.022806
17	0.0000000	113	0.64228	0.96794	0.022806

[Hide](#)

```
rpart.plot(tree_bos3_cp0,
           main = "Classification Tree with CP = 0")
```



[Hide](#)

```
# results before pruning
cp0 <- printcp(tree_bos3_cp0)
```

Classification tree:

```
rpart(formula = cancellation_policy ~ ., data = bos3_train, method = "class",
      xval = 5, cp = 0)
```

Variables actually used in tree construction:

```
[1] accommodates      bathrooms      bed_type      bedrooms      cleaning_fee
[6] instant_bookable  log_price     number_of_reviews  review_scores_rating room_type
```

Root node error: 998/2083 = 0.47912

n= 2083

	CP	nsplit	rel error	xerror	xstd
1	0.1272545	0	1.00000	1.00000	0.022846
2	0.0115230	1	0.87275	0.87275	0.022557
3	0.0100200	3	0.84970	0.86874	0.022542
4	0.0046760	4	0.83968	0.85872	0.022504
5	0.0035070	7	0.82565	0.87375	0.022561
6	0.0027555	9	0.81864	0.88878	0.022613
7	0.0025050	14	0.80461	0.89579	0.022635
8	0.0020040	29	0.75150	0.90581	0.022666
9	0.0017307	44	0.71944	0.91283	0.022685
10	0.0015030	60	0.68838	0.90982	0.022677
11	0.0013360	66	0.67936	0.90281	0.022657
12	0.0010020	72	0.67134	0.92285	0.022712
13	0.0006680	89	0.65331	0.93587	0.022744
14	0.0005010	98	0.64729	0.94890	0.022771
15	0.0002505	104	0.64429	0.95892	0.022790
16	0.0002004	108	0.64329	0.96794	0.022806
17	0.0000000	113	0.64228	0.96794	0.022806

Hide

```
class(cp0)
```

```
[1] "matrix"
```


Hide

```
cp0 <- data.frame(cp0)
which.min(cp0$error)
```

```
[1] 4
```

[Hide](#)

```
plotcp(tree_bos3_cp0)
```


[Hide](#)

```
tree_cp0_pred <- predict(tree_bos3_cp0, bos3_train, type = "class")
confusionMatrix(tree_cp0_pred, bos3_train$CancellationPolicy)
```

Confusion Matrix and Statistics

	Reference		
Prediction	flexible	moderate	strict
flexible	325	68	106
moderate	45	246	108
strict	137	177	871

Overall Statistics

Accuracy : 0.6923
 95% CI : (0.6719, 0.712)
 No Information Rate : 0.5209
 P-Value [Acc > NIR] : < 0.00000000000000022

Kappa : 0.4873

Mcnemar's Test P-Value : 0.0000131

Statistics by Class:

	Class: flexible	Class: moderate	Class: strict
Sensitivity	0.6410	0.5010	0.8028
Specificity	0.8896	0.9039	0.6854
Pos Pred Value	0.6513	0.6165	0.7350
Neg Pred Value	0.8851	0.8545	0.7617
Prevalence	0.2434	0.2357	0.5209
Detection Rate	0.1560	0.1181	0.4181
Detection Prevalence	0.2396	0.1916	0.5689
Balanced Accuracy	0.7653	0.7025	0.7441

[Hide](#)

```
tree_cp0_pred2 <- predict(tree_bos3_cp0, bos3_valid, type = "class")
confusionMatrix(tree_cp0_pred2, bos3_valid$cancellation_policy)
```

Confusion Matrix and Statistics

	Reference		
Prediction	flexible	moderate	strict
flexible	170	72	106
moderate	55	81	131
strict	112	173	485

Overall Statistics

Accuracy : 0.5314
 95% CI : (0.5047, 0.558)
 No Information Rate : 0.5213
 P-Value [Acc > NIR] : 0.23393

Kappa : 0.2238

Mcnemar's Test P-Value : 0.04124

Statistics by Class:

	Class: flexible	Class: moderate	Class: strict
Sensitivity	0.5045	0.24847	0.6717
Specificity	0.8302	0.82436	0.5701
Pos Pred Value	0.4885	0.30337	0.6299
Neg Pred Value	0.8390	0.78086	0.6146
Prevalence	0.2433	0.23538	0.5213
Detection Rate	0.1227	0.05848	0.3502
Detection Prevalence	0.2513	0.19278	0.5560
Balanced Accuracy	0.6673	0.53641	0.6209

[Hide](#)

```
# rpart with xerror minimizing cp value
tree_bos3_cp_min_error <- rpart(cancellation_policy ~ .,
                                data=bos3_train,
                                method = "class",
                                cp = 0.0050100) # complexity parameter that corresponds to the value where the xerror was minimized (4th record in the cp0 df, nsplit = 7)
printcp(tree_bos3_cp_min_error)
```

Classification tree:

```
rpart(formula = cancellation_policy ~ ., data = bos3_train, method = "class",
      cp = 0.00501)
```

Variables actually used in tree construction:

```
[1] cleaning_fee      log_price          number_of_reviews  review_scores_rating
```

Root node error: 998/2083 = 0.47912

n= 2083

	CP	nsplit	rel error	xerror	xstd
1	0.127255	0	1.00000	1.00000	0.022846
2	0.011523	1	0.87275	0.87275	0.022557
3	0.010020	3	0.84970	0.87074	0.022550
4	0.005010	4	0.83968	0.85872	0.022504

[Hide](#)

```
# results after pruning
cp_min_error.pred <- predict(tree_bos3_cp_min_error, bos3_train, type = "class")
confusionMatrix(cp_min_error.pred, bos3_train$cancellation_policy)
```

Confusion Matrix and Statistics

	Reference		
Prediction	flexible	moderate	strict
flexible	244	71	93
moderate	6	16	7
strict	257	404	985

Overall Statistics

Accuracy : 0.5977
 95% CI : (0.5763, 0.6188)
 No Information Rate : 0.5209
 P-Value [Acc > NIR] : 0.0000000000001078

Kappa : 0.2514

Mcnemar's Test P-Value : < 0.00000000000000022

Statistics by Class:

	Class: flexible	Class: moderate	Class: strict
Sensitivity	0.4813	0.032587	0.9078
Specificity	0.8959	0.991834	0.3377
Pos Pred Value	0.5980	0.551724	0.5984
Neg Pred Value	0.8430	0.768744	0.7712
Prevalence	0.2434	0.235718	0.5209
Detection Rate	0.1171	0.007681	0.4729
Detection Prevalence	0.1959	0.013922	0.7902
Balanced Accuracy	0.6886	0.512210	0.6228

[Hide](#)

```
cp_min_error_pred2 <- predict(tree_bos3_cp_min_error, bos3_valid, type = "class")
confusionMatrix(cp_min_error_pred2, bos3_valid$cancellation_policy)
```

Confusion Matrix and Statistics

	Reference		
Prediction	flexible	moderate	strict
flexible	160	47	70
moderate	9	10	3
strict	168	269	649

Overall Statistics

Accuracy : 0.5913
 95% CI : (0.5649, 0.6174)
 No Information Rate : 0.5213
 P-Value [Acc > NIR] : 0.00000009496

Kappa : 0.2416

McNemar's Test P-Value : < 0.00000000000000022

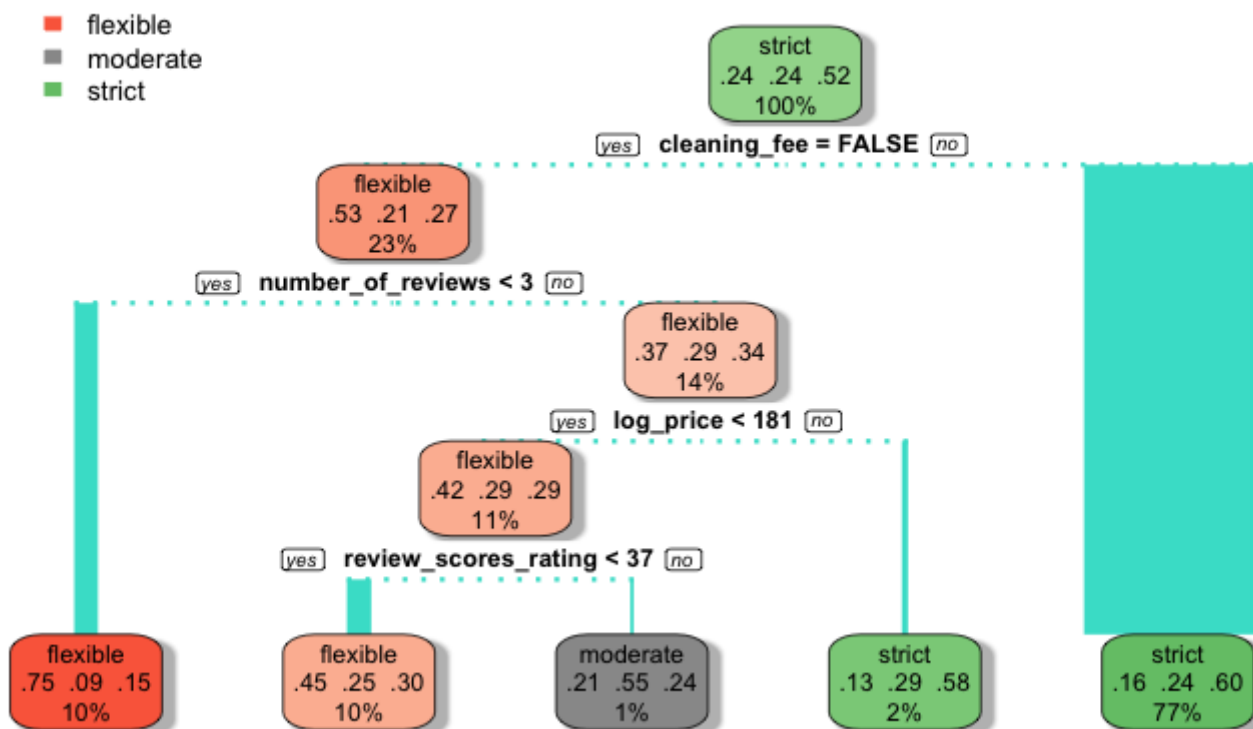
Statistics by Class:

	Class: flexible	Class: moderate	Class: strict
Sensitivity	0.4748	0.03067	0.8989
Specificity	0.8884	0.98867	0.3409
Pos Pred Value	0.5776	0.45455	0.5976
Neg Pred Value	0.8403	0.76816	0.7559
Prevalence	0.2433	0.23538	0.5213
Detection Rate	0.1155	0.00722	0.4686
Detection Prevalence	0.2000	0.01588	0.7841
Balanced Accuracy	0.6816	0.50967	0.6199

[Hide](#)


```
# rpart.plot visualization of pruned tree
rpart.plot(tree_bos3_cp_min_error,
  main = "Classification Tree with CP = 0.0038076",
  clip.right.labs = FALSE,
  type = 2,
  branch = .75,
  yesno = 2,
  under = FALSE,
  cex.main = 2.5,
  shadow.col = "gray",
  branch.col = "turquoise",
  branch.lwd = 3,
  branch.lty = 3,
  branch.type = 5,
  gap = 0)
```

Classification Tree with CP = 0.003807



Step IV: Clustering

[Hide](#)

```
library(cluster)    # clustering algorithms
library(gridExtra)
```

Attaching package: 'gridExtra'

The following object is masked from 'package:dplyr':

combine

[Hide](#)

```
library(cluster)    # clustering algorithms
library(factoextra) # clustering algorithms & visualization
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

[Hide](#)

```
Data <- boston1
"Converted the categorical variable 'cancellation_policy' to numeric with a scale of 1 (flexible) to 5 (super_strict_60)."
```

```
[1] "Converted the categorical variable 'cancellation_policy' to numeric with a scale of 1 (flexible) to 5 (super_strict_60)."
```

[Hide](#)

```
levels(Data$cancellation_policy)
```

```
[1] "flexible"      "moderate"      "strict"        "super_strict_30" "super_strict_60"
```

Hide

```
# Converting cancellation_policy to numeric
Data$cancellation_policy<-revalue(Data$cancellation_policy,c("flexible"=1,"moderate"=2,"strict"=3,
                                                                "super_strict_30"=4, "super_strict_60"=5))

"Excluded the 'Cambridge' and 'Somerville' neighborhoods as there were only 5 properties in total."
```

```
[1] "Excluded the 'Cambridge' and 'Somerville' neighborhoods as there were only 5 properties in total."
```

Hide

```
Data <- filter(Data, neighbourhoood!= "Cambridge")
Data <- filter(Data, neighbourhoood!= "Somerville")
# Adding nightly price per person
Data <- Data%>%
  mutate(price_per_person = nightly_price/accommodates)
# Remove non-numeric columns
colnames(Data)
```

```
[1] "id"                "log_price"          "property_type"      "room_type"
[5] "amenities"         "accommodates"       "bathrooms"         "bed_type"
[9] "cancellation_policy" "cleaning_fee"       "description"        "host_has_profile_pic"
[13] "host_identity_verified" "host_response_rate" "host_since"         "instant_bookable"
[17] "latitude"          "longitude"          "name"              "neighbourhoood"
[21] "number_of_reviews" "review_scores_rating" "bedrooms"          "beds"
[25] "nightly_price"     "price_per_person"
```

Hide

```
Data<-Data[, -c(1,3:5,8,10:19)]
colnames(Data)
```


```
[1] "log_price"          "accommodates"       "bathrooms"          "cancellation_policy"
[5] "neighbourhoood"    "number_of_reviews"  "review_scores_rating" "bedrooms"
[9] "beds"              "nightly_price"      "price_per_person"
```

Hide


```
Aggregate_Data<-aggregate(cbind(log_price,accommodates,bathrooms,cancellation_policy, number_of_reviews,review_scores_rating,bedrooms,beds,nightly_price,price_per_person)~neighbourhood,data=Data,mean)

Boston<-data.frame(Aggregate_Data[,-1],row.names=Aggregate_Data$neighbourhood)
Scaled_Data <- scale(Boston)


# Optimal No. of Clusters
# elbow method
set.seed(123)
fviz_nbclust(Scaled_Data, kmeans, method = "wss")
```

[Hide](#)

```
# avg silhouette method
fviz_nbclust(Scaled_Data, kmeans, method = "silhouette")
```

[Hide](#)

```
# K-Means Algorithm
set.seed(123)
k4 <- kmeans(Scaled_Data, 4, nstart = 25)
p1 <- fviz_cluster(k4, data = Scaled_Data)
p1
```



Hide

```
print(k4)
```

K-means clustering with 4 clusters of sizes 3, 10, 8, 11

Cluster means:

	log_price	accommodates	bathrooms	cancellation_policy	number_of_reviews	review_scores_rating	bedrooms
1	-1.2777191	-1.8855414	-1.1909272	-1.7572116	-1.2601805	1.0808591	-1.1637506
2	0.7411959	0.1455015	-0.6815972	0.4960988	-0.1633199	-0.3921439	-0.9893809
3	0.8740255	1.0619844	1.1303778	0.2483884	0.2006509	0.3117588	1.0584553
4	-0.9610005	-0.3903878	0.1223393	-0.1524055	0.3462304	-0.1650190	0.4470380

	beds	nightly_price	price_per_person
1	-1.78940736	-1.3474056	-0.8553313
2	-0.07369644	0.5835037	0.6883243
3	0.84539474	1.0538433	0.7997643
4	-0.05981559	-0.9294152	-0.9741239

Clustering vector:

Allston-Brighton	Back Bay	Beacon Hill	Brookline	Charlestown
4	2	2	4	3
Chestnut Hill	Chinatown	Coolidge Corner	Dorchester	Downtown
1	3	2	4	3
Downtown Crossing	East Boston	Fenway/Kenmore	Financial District	Government Center
2	4	2	3	2
Harvard Square	Hyde Park	Jamaica Plain	Leather District	Mattapan
2	4	4	3	4
Mission Hill	Newton	North End	Roslindale	Roxbury
4	1	3	4	4
South Boston	South End	Theater District	Watertown	West End
3	2	2	2	3
West Roxbury	Winthrop			
4	1			

Within cluster sum of squares by cluster:

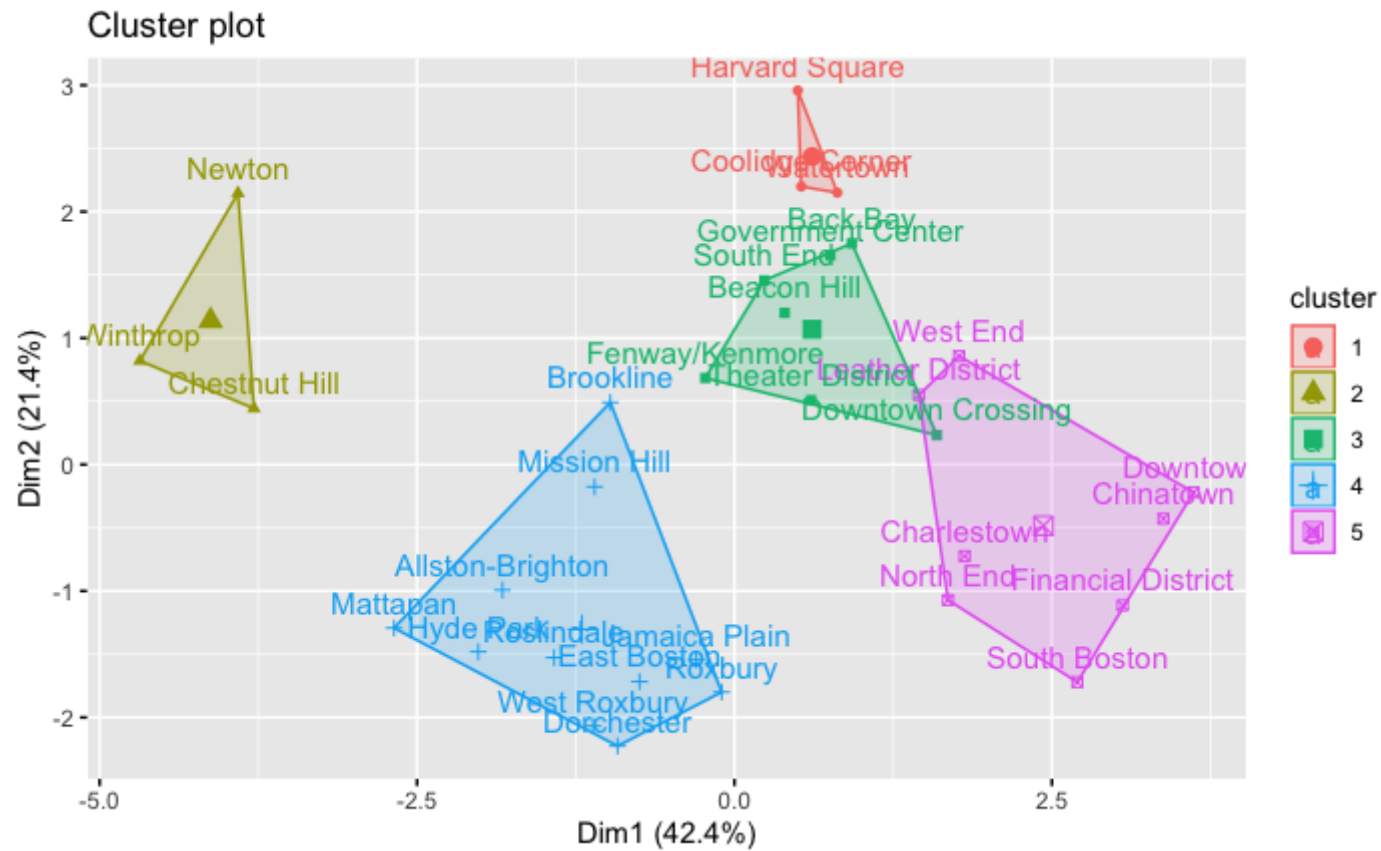
```
[1] 4.919815 52.736630 31.352756 38.045305
(between_SS / total_SS = 59.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"    "size"
[8] "iter"         "ifault"
```

Hide

```
k5 <- kmeans(Scaled_Data, 5, nstart = 25)
p2 <- fviz_cluster(k5, data = Scaled_Data)
p2
```


[Hide](#)

```
print(k5)
```


K-means clustering with 5 clusters of sizes 3, 3, 7, 11, 8

Cluster means:

	log_price	accommodates	bathrooms	cancellation_policy	number_of_reviews	review_scores_rating	bedrooms
1	1.0628585	0.5436913	-1.3573021	0.09577922	-1.3957647	0.9853532	-1.3922558
2	-1.2777191	-1.8855414	-1.1909272	-1.75721157	-1.2601805	1.0808591	-1.1637506
3	0.6033405	-0.0251513	-0.3920095	0.66766432	0.3648707	-0.9824998	-0.8167202
4	-0.9610005	-0.3903878	0.1223393	-0.15240548	0.3462304	-0.1650190	0.4470380
5	0.8740255	1.0619844	1.1303778	0.24838838	0.2006509	0.3117588	1.0584553

	beds	nightly_price	price_per_person
1	0.90633476	0.7781191	0.4305877
2	-1.78940736	-1.3474056	-0.8553313
3	-0.49370981	0.5000971	0.7987828
4	-0.05981559	-0.9294152	-0.9741239
5	0.84539474	1.0538433	0.7997643

Clustering vector:

Allston-Brighton	Back Bay	Beacon Hill	Brookline	Charlestown
4	3	3	4	5
Chestnut Hill	Chinatown	Coolidge Corner	Dorchester	Downtown
2	5	1	4	5
Downtown Crossing	East Boston	Fenway/Kenmore	Financial District	Government Center
3	4	3	5	3
Harvard Square	Hyde Park	Jamaica Plain	Leather District	Mattapan
1	4	4	5	4
Mission Hill	Newton	North End	Roslindale	Roxbury
4	2	5	4	4
South Boston	South End	Theater District	Watertown	West End
5	3	3	1	5
West Roxbury	Winthrop			
4	2			

Within cluster sum of squares by cluster:


```
[1] 8.960427 4.919815 20.109007 38.045305 31.352756
(between_SS / total_SS = 66.6 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"    "size"
[8] "iter"         "ifault"
```

Hide

```
k6 <- kmeans(Scaled_Data, 6, nstart = 25)
p3 <- fviz_cluster(k6, data = Scaled_Data)
p3
```



Hide

```
print(k6)
```

K-means clustering with 6 clusters of sizes 7, 1, 10, 3, 8, 3

Cluster means:

	log_price	accommodates	bathrooms	cancellation_policy	number_of_reviews	review_scores_rating	bedrooms
1	0.6033405	-0.0251513	-0.3920095	0.66766432	0.3648707	-0.9824998	-0.8167202
2	-0.2878330	0.2664419	-1.3573021	1.86454316	-1.3957647	0.9853532	1.0070489
3	-1.0283173	-0.4560708	0.2703034	-0.35410034	0.5204299	-0.2800562	0.3910369
4	-1.2777191	-1.8855414	-1.1909272	-1.75721157	-1.2601805	1.0808591	-1.1637506
5	0.8740255	1.0619844	1.1303778	0.24838838	0.2006509	0.3117588	1.0584553
6	1.0628585	0.5436913	-1.3573021	0.09577922	-1.3957647	0.9853532	-1.3922558

	beds	nightly_price	price_per_person
1	-0.49370981	0.5000971	0.7987828
2	-0.90038602	-0.7399561	-1.2765084
3	0.02424146	-0.9483611	-0.9438855
4	-1.78940736	-1.3474056	-0.8553313
5	0.84539474	1.0538433	0.7997643
6	0.90633476	0.7781191	0.4305877

Clustering vector:

Allston-Brighton	Back Bay	Beacon Hill	Brookline	Charlestown
3	1	1	2	5
Chestnut Hill	Chinatown	Coolidge Corner	Dorchester	Downtown
4	5	6	3	5
Downtown Crossing	East Boston	Fenway/Kenmore	Financial District	Government Center
1	3	1	5	1
Harvard Square	Hyde Park	Jamaica Plain	Leather District	Mattapan
6	3	3	5	3
Mission Hill	Newton	North End	Roslindale	Roxbury
3	4	5	3	3
South Boston	South End	Theater District	Watertown	West End
5	1	1	6	5
West Roxbury	Winthrop			
3	4			

Within cluster sum of squares by cluster:

```
[1] 20.109007 0.000000 24.133160 4.919815 31.352756 8.960427
(between_SS / total_SS = 71.1 %)
```


Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"    "size"
```

```
[8] "iter"         "ifault"
```

Hide

```
k7 <- kmeans(Scaled_Data, 7, nstart = 25)
p4 <- fviz_cluster(k7, data = Scaled_Data)
p4
```



Hide

```
print(k7)
```

K-means clustering with 7 clusters of sizes 3, 3, 7, 1, 5, 10, 3

Cluster means:

	log_price	accommodates	bathrooms	cancellation_policy	number_of_reviews	review_scores_rating	bedrooms
1	1.0628585	0.5436913	-1.35730206	0.09577922	-1.39576469	0.9853532	-1.3922558
2	-1.2777191	-1.8855414	-1.19092717	-1.75721157	-1.26018046	1.0808591	-1.1637506
3	0.8589479	1.1774613	1.10181853	0.18661073	0.39085095	0.4153587	1.1659898
4	-0.2878330	0.2664419	-1.35730206	1.86454316	-1.39576469	0.9853532	1.0070489
5	0.6112883	-0.2449134	0.06828808	0.37879885	-0.02660866	-0.3054778	-0.5410740
6	-1.0283173	-0.4560708	0.27030339	-0.35410034	0.52042986	-0.2800562	0.3910369
7	0.7155037	0.4340513	-0.58507142	1.15349597	0.51879607	-1.9211835	-0.9019859

	beds	nightly_price	price_per_person
1	0.90633476	0.7781191	0.4305877
2	-1.78940736	-1.3474056	-0.8553313
3	1.02728813	1.0555134	0.7320034
4	-0.90038602	-0.7399561	-1.2765084
5	-0.67760477	0.5984635	1.1614178
6	0.02424146	-0.9483611	-0.9438855
7	-0.16526795	0.5168383	0.3528272

Clustering vector:

Allston-Brighton	Back Bay	Beacon Hill	Brookline	Charlestown
6	5	5	4	3
Chestnut Hill	Chinatown	Coolidge Corner	Dorchester	Downtown
2	3	1	6	3
Downtown Crossing	East Boston	Fenway/Kenmore	Financial District	Government Center
7	6	5	3	7
Harvard Square	Hyde Park	Jamaica Plain	Leather District	Mattapan
1	6	6	3	6
Mission Hill	Newton	North End	Roslindale	Roxbury
6	2	3	6	6
South Boston	South End	Theater District	Watertown	West End
3	5	7	1	5
West Roxbury	Winthrop			
6	2			

Within cluster sum of squares by cluster:

```
[1] 8.960427 4.919815 24.949347 0.000000 9.446952 24.133160 7.013507
(between_SS / total_SS = 74.4 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss" "betweenss"    "size"
[8] "iter"        "ifault"
```

[Hide](#)

```
# comparing the clusters
# grid.arrange(p1, p2, p3, p4, nrow = 2)

# Adding to our initial data to do some descriptive statistics at the cluster level
k4$centers
```

```
  log_price accommodates  bathrooms cancellation_policy number_of_reviews review_scores_rating  bedrooms
1 -1.2777191  -1.8855414 -1.1909272         -1.7572116         -1.2601805          1.0808591 -1.1637506
2  0.7411959   0.1455015 -0.6815972          0.4960988         -0.1633199         -0.3921439 -0.9893809
3  0.8740255   1.0619844  1.1303778          0.2483884          0.2006509          0.3117588  1.0584553
4 -0.9610005  -0.3903878  0.1223393         -0.1524055          0.3462304         -0.1650190  0.4470380

  beds nightly_price price_per_person
1 -1.78940736   -1.3474056       -0.8553313
2 -0.07369644    0.5835037        0.6883243
3  0.84539474    1.0538433        0.7997643
4 -0.05981559   -0.9294152       -0.9741239
```

[Hide](#)

```
Boston %>%
  mutate(Cluster = k4$cluster) %>%
  group_by(Cluster) %>%
  summarise_all("mean")
```

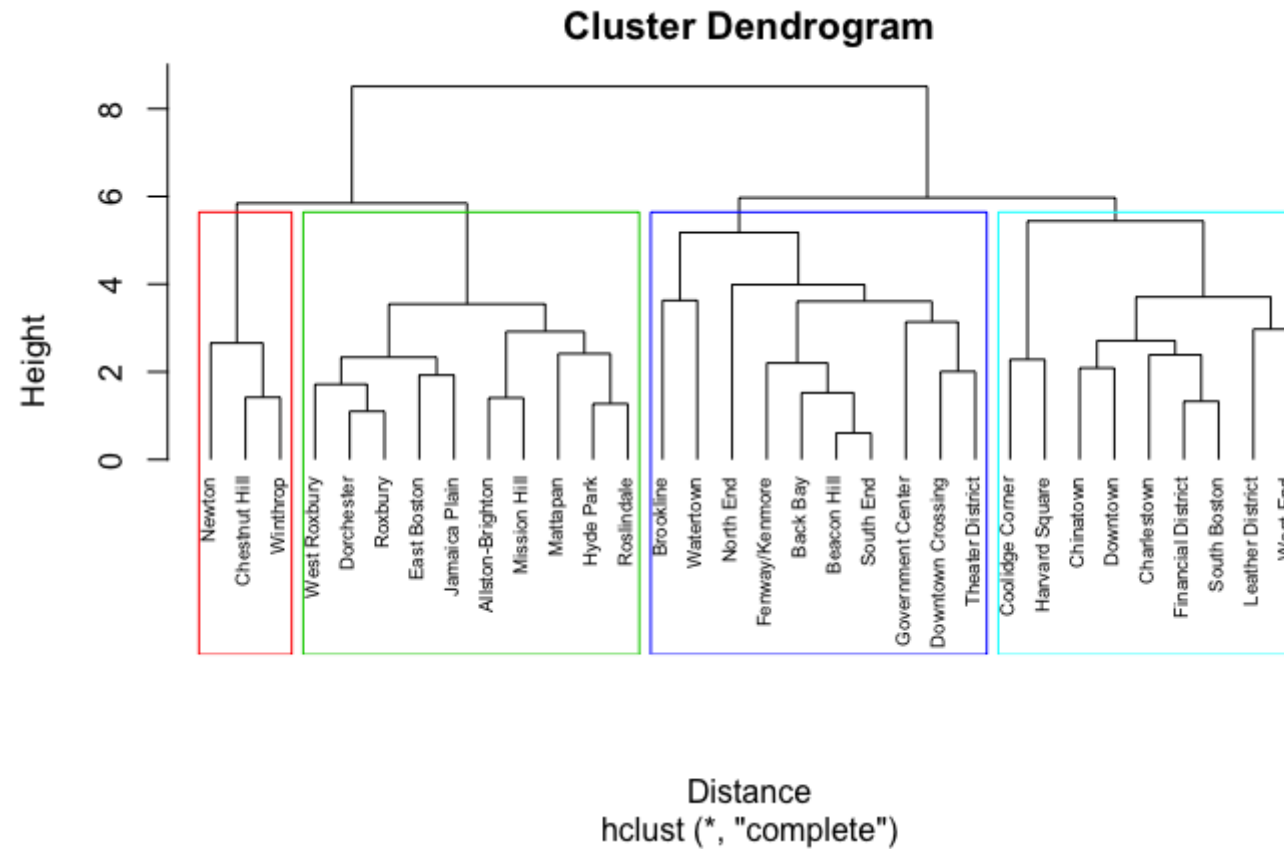
Cluster <int>	log_price <dbl>	accommodates <dbl>	bathrooms <dbl>	cancellation_policy <dbl>	number_of_reviews <dbl>
1	4.376402	2.206349	1.023810	1.634921	1.904762
2	5.256466	3.427297	1.096699	2.484218	17.314063
3	5.314367	3.978235	1.356006	2.390853	22.427325

Cluster <int>	log_price <dbl>	accommodates <dbl>	bathrooms <dbl>	cancellation_policy <dbl>	number_of_reviews <dbl>
4	4.514463	3.105151	1.211748	2.239790	24.472506

4 rows | 1-6 of 11 columns

[Hide](#)

```
# Hierarchical Clustering
# Dissimilarity matrix
Distance <- dist(Scaled_Data, method = "euclidean")
# Hierarchical clustering using Complete Linkage
hc <- hclust(Distance, method = "complete" )
# Plot the obtained dendrogram
plot(hc, cex = 0.6, hang = -1)
rect.hclust(hc, k = 4, border = 2:5)
```



Step V: Conclusions

[Code](#)

[1] "The overall process was very collaborative in nature. As a team we had decided to prepare and explore the data together and then work on our individual areas to come up with our analysis. In the process of doing so we had discussed and commented constructively on each other's work which resulted in a much better quality output in the end. The exploratory analysis of the Airbnb data helped us understand the rental landscape of Boston through various statistics and visualizations. For e.g. the clustering analysis shows how certain neighborhoods are similar in nature and also what are the various characteristics that make them a part of each cluster. This helps individuals and businesses alike to answer certain questions like which neighborhoods have the highest review scores, listing price etc. \n\nThe classification tree model could be useful for a property owner who is interested in listing their property for rental on Airbnb. If they were not sure what type of cancellation policy they should implement for their new rental, Airbnb could provide a service that helped them with "Based on the characteristics that you have provided regarding your potential listing, and other properties that share some of these characteristics, we recommend a cancellation policy of "x"." If you operate under the presumption that the existing properties from which the model was built have their cancellation policies for good reason, this will help the new owner arrive at a good decision from the get-go. The property owner, Airbnb, and even potentially the ultimate customer/renter can all benefit from implementation of a Classification Tree such as the one we constructed. \n\nWhen we examine the Naive Bayes model, we know that the model with its selected variables is pretty accurate. Upon setting up the categorical bins derived from the log price variable, we can see that the majority of Airbnb rentals were conducted in the below average price category. This plays an important role in realizing how much visitors are willing to pay to stay in Boston. With the mean and median log price falling at 4.913/4.884 we can see that the majority would rather pay less than the median/mean log price. In its true dollar format, this comes out to be around \$80 a night but less than \$136 dollars. \n\nAn important question to consider is who are these visitors? By using Naive Bayes alone it is difficult to determine who these visitors are. However, when you combine our Naive Bayes model along with other models such as the K-Means Analysis, you will be able to see which neighborhoods the 'below average price category' would be placed in and depict a clearer picture of who the audience are and their needs. Boston is known for being a college town with most colleges centered at or around the city. If there is anyone who is willing to pay a below average price rating, it would most likely be college students which can be seen as the majority of would fall at or below average price category. But you also have another half that is unknown and that half for certain falls in above average category.\n\nWe can therefore conclude that the Airbnb data does give an insight as to who is visiting Boston and what is the price that they are willing to pay to stay. By using variables such as the price, cancellation policy, cleaning fee, neighborhoods, etc., we can derive further in depth questions as to the demographics of renting in Boston.\n"